

# Linear supervised regression

---

## 0. Import library

---

Import library

In [4]:

```
# Import libraries

# math library
import numpy as np

# visualization library
%matplotlib inline
from IPython.display import set_matplotlib_formats
set_matplotlib_formats('png2x', 'pdf')
import matplotlib.pyplot as plt

# machine learning library
from sklearn.linear_model import LinearRegression

# 3d visualization
from mpl_toolkits.mplot3d import axes3d

# computational time
import time
```

## 1. Load dataset

---

Load a set of data pairs  $\{x_i, y_i\}_{i=1}^n$  where  $x$  represents label and  $y$  represents target.

In [5]:

```
# import data with numpy
data = np.loadtxt('profit_population.txt', delimiter=',')
```

## 2. Explore the dataset distribution

---

Plot the training data points.

In [6]:

```


x_train = data[:,0]
y_train = data[:,1]

plt.scatter(x_train,y_train,c="r",marker="x")
plt.xlabel("popluation size(x 10$)")
plt.ylabel("profit(x 10k$)")
plt.title("Training data")
plt.plot()

```

Out[6]:

[]



### 3. Define the linear prediction function

$$f_w(x) = w_0 + w_1x$$

**Vectorized implementation:**

$$f_w(x) = Xw$$

with

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \quad \Rightarrow \quad f_w(x) = Xw = \begin{bmatrix} w_0 + w_1x_1 \\ w_0 + w_1x_2 \\ \vdots \\ w_0 + w_1x_n \end{bmatrix}$$

Implement the vectorized version of the linear predictive function.

In [7]:

```

# construct data matrix
X = np.array(x_train).reshape(-1,1)
X = np.concatenate((np.ones(X.shape),X),axis=1)
# parameters vector
w = np.array([[1],[1]])

# predictive function definition
def f_pred(X,w):

    f = np.matmul(X,w)

    return f

# Test predictive function
y_pred = f_pred(X,w)
y_pred.shape

```

Out[7]:

(97, 1)

## 4. Define the linear regression loss

$$L(w) = \frac{1}{n} \sum_{i=1}^n \left( f_w(x_i) - y_i \right)^2$$

**Vectorized implementation:**

$$L(w) = \frac{1}{n} (Xw - y)^T (Xw - y)$$

with

$$Xw = \begin{bmatrix} w_0 + w_1 x_1 \\ w_0 + w_1 x_2 \\ \vdots \\ w_0 + w_1 x_n \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Implement the vectorized version of the linear regression loss function.

In [8]:

```
# loss function definition
def loss_mse(y_pred,y):
    temp=y_pred-y
    loss = np.matmul(temp.T,temp)/temp.size

    return loss.item()

# Test loss function
y = np.array(y_train).reshape(-1,1)# label
y_pred = f_pred(X,w)# prediction

loss = loss_mse(y_pred,y)
```

## 5. Define the gradient of the linear regression loss

**Vectorized implementation: Given the loss**

$$L(w) = \frac{1}{n}(Xw - y)^T(Xw - y)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T (Xw - y)$$

Implement the vectorized version of the gradient of the linear regression loss function.

In [9]:

```
# gradient function definition
def grad_loss(y_pred,y,X):
    grad = 2*np.matmul(X.T,y_pred-y)/y_pred.size
    return grad

# Test grad function
y_pred = f_pred(X,w)
grad = grad_loss(y_pred,y,X)
```

## 6. Implement the gradient descent algorithm

- Vectorized implementation:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T (Xw^k - y)$$

**Implement the vectorized version of the gradient descent function.**

**Plot the loss values  $L(w^k)$  with respect to iteration  $k$  the number of iterations.**

In [10]:

```
# gradient descent function definition
def grad_desc(X, y, w_init, tau, max_iter):

    L_iters = np.empty(max_iter) # record the loss values
    w_iters = np.empty([max_iter, 2]) # record the parameter values
    w = w_init # initialization

    for i in range(max_iter): # loop over the iterations

        y_pred = f_pred(X, w) # linear prediction function
        grad_f = grad_loss(y_pred, y, X) # gradient of the loss
        w = w - tau * grad_f # update rule of gradient descent
        L_iters[i] = loss_mse(y_pred, y) # save the current loss value
        w_iters[i, :] = w.reshape(2) # save the current w value

    return w, L_iters, w_iters


# run gradient descent algorithm
start = time.time()
w_init = np.array([[1], [1]])
tau = 0.01
max_iter = 1000

w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)

print('Time=', time.time() - start) # plot the computational cost
print(L_iters[max_iter-1]) # plot the last value of the loss
print(w) # plot the last value of the parameter w

# plot
plt.figure(2)
plt.plot(L_iters) # plot the loss curve
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.show()
```

```
Time= 0.02196812629699707  
8.957109868355534  
[[-3.76436927]  
 [ 1.17983192]]
```




## 7. Plot the linear prediction function

$$f_w(x) = w_0 + w_1x$$

In [11]:

```
# linear regression model
x_pred = np.linspace(0,25,100) # define the domain of the prediction function
y_pred = w[0]+x_pred*w[1]# compute the prediction values within the given domain x_pred

# plot
plt.figure(3)
plt.scatter(x_train,y_train,c="r",marker="x")
plt.plot(x_pred, y_pred,label="graident descent optimization" )
plt.legend(loc='best')
plt.title('Training data')
plt.xlabel('Population size (x 10k)')
plt.ylabel('Profit $(x 10k)')
plt.show()
```



## 8. Comparison with Scikit-learn linear regression algorithm

### Compare with the Scikit-learn solution

In [12]:

```
# run linear regression with scikit-learn
start = time.time()
lin_reg_sklearn = LinearRegression()
lin_reg_sklearn.fit(np.reshape(x_train,(-1,1)),np.reshape(y_train,(-1,1))) # learn the model parameters
print('Time=',time.time() - start)

# compute loss value
w_sklearn = np.zeros([2,1])
w_sklearn[0,0] = lin_reg_sklearn.intercept_
w_sklearn[1,0] = lin_reg_sklearn.coef_

print(w_sklearn)

loss_sklearn = loss_mse(np.reshape(lin_reg_sklearn.predict(np.reshape(y_train,(-1,1))),(-1,1)) ,
np.reshape(y_train,(-1,1))) # compute the loss from the sklearn solution

print('loss sklearn=', loss_sklearn)
print('loss gradient descent=',L_iters[-1])

# plot
y_pred_sklearn = lin_reg_sklearn.predict(np.reshape(x_pred,(-1,1)))# prediction obtained by the
sklearn library

plt.figure(3)

plt.scatter(x_train,y_train,c="r",marker="x")
plt.plot(x_pred, y_pred,label="gradient descent optimization" )
plt.plot(x_pred,y_pred_sklearn,label="Scikit-learn optimization" )
plt.legend(loc='best')
plt.title('Training data')
plt.xlabel('Population size (x 10k)')
plt.ylabel('Profit $(x 10k)')
plt.show()
```




Time= 0.020911216735839844

[[-3.89578088]

[ 1.19303364]]

loss sklearn= 8.785041210245968

loss gradient descent= 8.957109868355534



## 9. Plot the loss surface, the contours of the loss and the gradient descent steps

---

In [114]:

```

# plot gradient descent
def plot_gradient_descent(X,y,w_init,tau,max_iter):

    def f_pred(X,w):

        f = f = np.matmul(X,w)

        return f

    def loss_mse(y_pred,y):

        temp=y_pred-y
        loss = np.matmul(temp.T,temp)/temp.size

        return loss.item()

    # gradient descent function definition
    def grad_desc(X, y, w_init, tau, max_iter):

        L_iters = np.empty(max_iter)# record the loss values
        w_iters = np.empty([max_iter,2])# record the parameter values
        w = w_init # initialization

        for i in range(max_iter): # loop over the iterations

            y_pred = f_pred(X,w)# linear prediction function
            grad_f = grad_loss(y_pred,y,X)# gradient of the loss
            w = w-tau*grad_f# update rule of gradient descent
            L_iters[i] = loss_mse(y_pred,y)# save the current loss value
            w_iters[i,:] = w.reshape(2)# save the current w value

        return w, L_iters, w_iters

    # run gradient descent
    w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)

    # Create grid coordinates for plotting a range of L(w0,w1)-values
    B0 = np.linspace(-10, 10, 50)
    B1 = np.linspace(-1, 4, 50)

    xx, yy = np.meshgrid(B0, B1, indexing='xy')
    Z = np.zeros((B0.size,B1.size))

    # Calculate loss values based on L(w0,w1)-values
    for (i,j),v in np.ndenumerate(Z):
        w_0=i*0.4-10
        w_1=j*0.1-1
        w_temp=np.array([[w_0],[w_1]])
        y_pred_temp=f_pred(X,w_temp)
        Z[i,j] = loss_mse(y_pred_temp,y )

    # 3D visualization
    fig = plt.figure(figsize=(15,6))
    ax1 = fig.add_subplot(121)
    ax2 = fig.add_subplot(122, projection='3d')

    # Left plot

```

```

CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
ax1.scatter(x=w_iters[0,0],y=w_iters[0,1],c="red")
ax1.plot(w_iters[:,0],w_iters[:,1],c="blue")
plt.plot

# Right plot
ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.jet)
ax2.set_zlabel('Loss  $L(w_0, w_1)$ ')
ax2.set_zlim(Z.min(), Z.max())

# plot gradient descent
Z2 = np.zeros([max_iter])

for i in range(max_iter):
    w0 = w_iters[i][0]
    w1 = w_iters[i][1]
    w_temp=np.array([[w0],[w1]])
    y_pred_temp = f_pred(X,w_temp)# linear prediction function
    Z2[i] = loss_mse(y_pred_temp,y)

ax2.plot(w_iters[:,0],w_iters[:,1],Z2 )
ax2.scatter(w_iters[0,0],w_iters[0,1],Z2[0] )

# settings common to both plots
for ax in fig.axes:
    ax.set_xlabel(r'$w_0$', fontsize=17)
    ax.set_ylabel(r'$w_1$', fontsize=17)

```



In [118]:

```

# run plot_gradient_descent function
w_init = np.array([[ -5],[ 4]])
tau = 0.01
max_iter = 1000

plot_gradient_descent(X,y,w_init,tau,max_iter)

```



## Output results

### 1. Plot the training data (1pt)


In [46]:

```
data = np.loadtxt('profit_population.txt', delimiter=',')
x_train = data[:,0]
y_train = data[:,1]

plt.scatter(x_train,y_train,c="r",marker="x")
plt.xlabel("poplualtion size(x 10$)")
plt.ylabel("profit(x 10K$)")
plt.title("Training data")
plt.plot()
```

Out[46]:


[]



## 2. Plot the loss curve in the course of gradient descent (2pt)

In [120]:

```
# plot
plt.figure(2)
plt.plot(L_iters) # plot the loss curve
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.show()
```




### 3. Plot the prediction function superimposed on the training data (2pt)

In [122]:

```
x_pred = np.linspace(0,25,100)
y_pred = w[0]+x_pred*w[1]

plt.figure(3)
plt.scatter(x_train,y_train,c="r",marker="x")
plt.plot(x_pred, y_pred,label="graident descent optimization" )
plt.legend(loc='best')
plt.title('Training data')
plt.xlabel('Population size (x 10k)')
plt.ylabel('Profit $(x 10k)')
plt.show()
```



**4. Plot the prediction functions obtained by both the Scikit-learn linear regression solution and the gradient descent superimposed on the training data (2pt)**

In [123]:

```
lin_reg_sklearn = LinearRegression()
lin_reg_sklearn.fit(np.reshape(x_train,(-1,1)),np.reshape(y_train,(-1,1))) # learn the model parameters


w_sklearn = np.zeros([2,1])
w_sklearn[0,0] = lin_reg_sklearn.intercept_
w_sklearn[1,0] = lin_reg_sklearn.coef_

loss_sklearn = loss_mse(np.reshape(lin_reg_sklearn.predict(np.reshape(y_train,(-1,1))),(-1,1)) ,
np.reshape(y_train,(-1,1))) # compute the loss from the sklearn solution

y_pred_sklearn = lin_reg_sklearn.predict(np.reshape(x_pred,(-1,1)))# prediction obtained by the
sklearn library

plt.figure(3)

plt.scatter(x_train,y_train,c="r",marker="x")
plt.plot(x_pred, y_pred,label="gradient descent optimization" )
plt.plot(x_pred,y_pred_sklearn,label="Scikit-learn optimization" )
plt.legend(loc='best')
plt.title('Training data')
plt.xlabel('Population size (x 10k)')
plt.ylabel('Profit $(x 10k)')
plt.show()
```



## 5. Plot the loss surface (right) and the path of the gradient descent (2pt)

In [148]:

```

w_init=np.array([[ -3],[4]])
max_iter=1000
tau=0.01
w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)

B0 = np.linspace(-10, 10, 50)
B1 = np.linspace(-1, 4, 50)

xx, yy = np.meshgrid(B0, B1, indexing='xy')
Z = np.zeros((B0.size,B1.size))

for (i,j),v in np.ndenumerate(Z):
    w_0=i*0.4-10
    w_1=j*0.1-1
    w_temp=np.array([[w_0],[w_1]])
    y_pred_temp=f_pred(X,w_temp)
    Z[i,j] = loss_mse(y_pred_temp,y )

fig = plt.figure(figsize=(15,6))
ax2 = fig.add_subplot(122, projection='3d')

ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.jet)
ax2.set_zlabel('Loss  $L(w_0,w_1)$ ')
ax2.set_zlim(Z.min(),Z.max())

Z2 = np.zeros([max_iter])

for i in range(max_iter):
    w0 = w_iters[i][0]
    w1 = w_iters[i][1]
    w_temp=np.array([[w0],[w1]])
    y_pred_temp = f_pred(X,w_temp)# linear prediction function
    Z2[i] = loss_mse(y_pred_temp,y)

ax2.plot(w_iters[:,0],w_iters[:,1],Z2 )
ax2.scatter(w_iters[0,0],w_iters[0,1],Z2[0] )

ax2.set_xlabel(r' $w_0$ ', fontsize=17)
ax2.set_ylabel(r' $w_1$ ', fontsize=17)


plt.plot

```



Out [148]:

```
<function matplotlib.pyplot.plot(*args, scalex=True, scaley=True, data=None, **kwargs)>
```



**6. Plot the contour of the loss surface (left) and the path of the gradient descent (2pt)**

In [149]:

```

w_init=np.array([[ -3],[ 4]])
max_iter=1000
tau=0.01
w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)

B0 = np.linspace(-10, 10, 50)
B1 = np.linspace(-1, 4, 50)

xx, yy = np.meshgrid(B0, B1, indexing='xy')
Z = np.zeros((B0.size,B1.size))

for (i,j),v in np.ndenumerate(Z):
    w_0=i*0.4-10
    w_1=j*0.1-1
    w_temp=np.array([[w_0],[w_1]])
    y_pred_temp=f_pred(X,w_temp)
    Z[i,j] = loss_mse(y_pred_temp,y )

fig = plt.figure(figsize=(15,6))
ax1 = fig.add_subplot(121)

CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
ax1.scatter(x=w_iters[0,0],y=w_iters[0,1],c="red")
ax1.plot(w_iters[:,0],w_iters[:,1],c="blue")

ax1.set_xlabel(r'$w_0$', fontsize=17)
ax1.set_ylabel(r'$w_1$', fontsize=17)

plt.plot()

```

Out [149]:

[]

