

Supervised Logistic Regression for Classification

0. Import library

In [26]:

```
# Import libraries

# math library
import numpy as np

# visualization library
%matplotlib inline
from IPython.display import set_matplotlib_formats
set_matplotlib_formats('png2x', 'pdf')
import matplotlib.pyplot as plt

# machine learning library
from sklearn.linear_model import LogisticRegression

# 3d visualization
from mpl_toolkits.mplot3d import axes3d

# computational time
import time
```

1. Load dataset

The data features $x_i = (x_{i(1)}, x_{i(2)})$ represent 2 exam grades $x_{i(1)}$ and $x_{i(2)}$ for each student i .

The data label y_i indicates if the student i was admitted (value is 1) or rejected (value is 0).

In [27]:

```
# import data with numpy
data = np.loadtxt('dataset.txt', delimiter=',')

# number of training data
n = data.shape[0]
print('Number of training data=', n)
```

Number of training data= 100

2. Explore the dataset distribution


Plot the training data points.

You may use matplotlib function `scatter(x,y)`.

In [28]:

```
x1 = data[:,0] # exam grade 1
x2 = data[:,1] # exam grade 2
idx_admit = (data[:,2]==1) # index of students who were admitted
idx_rejec = (data[:,2]==0) # index of students who were rejected

plt.figure(figsize=(8,8))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )
plt.title('Training data')
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Admitted", "Rejected"])
plt.show()
```



In [29]:

x2

Out [29]:

```
array([78.02469282, 43.89499752, 72.90219803, 86.3085521 , 75.34437644,
       56.31637178, 96.51142588, 46.55401354, 87.42056972, 43.53339331,
       38.22527806, 30.60326323, 76.4819633 , 97.71869196, 76.03681085,
       89.20735014, 52.74046973, 46.67857411, 92.92713789, 47.57596365,
       42.83843832, 65.79936593, 48.85581153, 44.2095286 , 68.97235999,
       69.95445795, 44.82162893, 38.80067034, 50.25610789, 64.99568096,
       72.80788731, 57.05198398, 63.12762377, 69.43286012, 71.16774802,
       52.21388588, 98.86943574, 80.90806059, 41.57341523, 75.23772034,
       56.30804622, 46.85629026, 65.56892161, 40.61825516, 45.82270146,
       52.06099195, 70.4582 , 86.72782233, 96.76882412, 88.69629255,
       74.16311935, 60.999031 , 43.39060181, 60.39634246, 49.80453881,
       59.80895099, 68.86157272, 95.59854761, 69.82457123, 78.45356225,
       85.75993667, 47.02051395, 39.26147251, 49.59297387, 66.45008615,
       41.09209808, 97.53518549, 51.88321182, 92.11606081, 60.99139403,
       43.30717306, 78.03168802, 96.22759297, 73.0949981 , 75.85844831,
       72.36925193, 88.475865 , 75.80985953, 42.50840944, 42.71987854,
       69.8037889 , 45.6943068 , 66.58935318, 59.51428198, 90.9601479 ,
       85.5943071 , 78.844786 , 90.424539 , 96.64742717, 60.76950526,
       77.15910509, 87.50879176, 35.57070347, 84.84513685, 45.35828361,
       48.3802858 , 87.10385094, 68.77540947, 64.93193801, 89.5298129 ])
```

3. Sigmoid/logistic function

$$\sigma(\eta) = \frac{1}{1 + \exp^{-\eta}}$$

Define and plot the sigmoid function for values in [-10,10]:


You may use functions `np.exp` , `np.linspace` .

In [30]:

```
def sigmoid(z):
    sigmoid_f = 1/(1+np.exp(-z))
    return sigmoid_f

# plot
x_values = np.linspace(-10,10)

plt.figure(2)
plt.plot(x_values,sigmoid(x_values))
plt.title("Sigmoid function")
plt.grid(True)
```



4. Define the prediction function for the classification

The prediction function is defined by:

$$p_w(x) = \sigma(w_0 + w_1x_{(1)} + w_2x_{(2)}) = \sigma(w^T x)$$

Implement the prediction function in a vectorised way as follows:

$$X = \begin{bmatrix} 1 & x_{1(1)} & x_{1(2)} \\ 1 & x_{2(1)} & x_{2(2)} \\ \vdots & \vdots & \vdots \\ 1 & x_{n(1)} & x_{n(2)} \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} \quad \Rightarrow \quad p_w(x) = \sigma(Xw) = \begin{bmatrix} \sigma(w_0 + w_1x_{1(1)} + w_2x_{1(2)}) \\ \sigma(w_0 + w_1x_{2(1)} + w_2x_{2(2)}) \\ \vdots \\ \sigma(w_0 + w_1x_{n(1)} + w_2x_{n(2)}) \end{bmatrix}$$

Use the new function `sigmoid`.

In [31]:

```
n = x1.size
X = np.stack([np.ones(n),x1,x2],axis=1)

# parameters vector
w = np.array([[0.3],[0],[0]])
w =np.array([-10,0.1,-0.2])[:,None]
# predictive function definition
def f_pred(X,w):

    p = sigmoid(np.dot(X,w).reshape(-1,1))

    return p
def f_pred2(X,w):

    f = np.matmul(X,w)

    return f
y_pred = f_pred(X,w)
y_pred
y_pred2=f_pred2(X,w)
y_pred2
```

Out[31]:

```
array([[ -22.1425726 ],
       [ -15.75032843],
       [ -20.99569873],
       [ -21.24345048],
       [ -17.16560168],
       [ -16.75494661],
       [ -23.19161872],
       [ -11.80832815],
       [ -19.87423527],
       [ -10.26339667],
       [  -8.0589001 ],
       [  -8.61928681],
       [ -17.06568732],
       [ -22.60727952],
       [ -21.25352826],
       [ -22.44436481],
       [ -13.64107954],
       [ -12.54102927],
       [ -21.51927662],
       [ -11.81731436],
       [ -11.83048491],
       [ -14.19219561],
       [ -14.71768348],
       [ -15.42069962],
       [ -16.00206285],
       [ -17.76379022],
       [ -10.94530771],
       [  -8.44869519],
       [ -13.86820098],
       [ -19.12055581],
       [ -18.42364852],
       [ -12.86994486],
       [ -17.41472678],
       [ -18.68203155],
       [ -20.20986023],
       [ -14.97926662],
       [ -26.38233714],
       [ -19.76391323],
       [ -10.83575775],
       [ -21.62918006],
       [ -12.87136988],
       [ -14.21648603],
       [ -13.66944754],
       [  -9.88677566],
       [ -14.05976511],
       [ -14.18993081],
       [ -16.37233651],
       [ -17.56840454],
       [ -23.14645845],
       [ -18.58276106],
       [ -16.83814208],
       [ -12.27255351],
       [  -9.62344895],
       [ -18.62681711],
       [ -14.93225815],
       [ -17.00312248],
       [ -14.00775115],
       [ -25.86198951],
       [ -16.54004511],
```

```
[-18.51106624],  
[-19.61242619],  
[-15.87549151],  
[-12.22691275],  
[-16.91271253],  
[-18.82319106],  
[-11.56233017],  
[-25.461282  ],  
[-15.46938604],  
[-20.39525476],  
[-15.52360695],  
[-15.38915131],  
[-19.20240556],  
[-22.01086917],  
[-18.57321105],  
[-19.28759404],  
[-14.49106461],  
[-22.96874609],  
[-20.11615593],  
[-12.45612626],  
[-10.32130955],  
[-15.06936814],  
[ -9.65541069],  
[-16.58594489],  
[-16.17898576],  
[-20.15535398],  
[-20.27200924],  
[-21.56141175],  
[-20.5371376  ],  
[-21.465943  ],  
[-16.91910065],  
[-16.02238791],  
[-18.45690326],  
[-11.56592458],  
[-19.51975813],  
[-10.08707605],  
[-11.32714088],  
[-23.19460011],  
[-13.82358101],  
[-17.45238585],  
[-20.42837328]])
```

In [32]:

```
X = np.stack([np.ones(n),x1,x2],axis=1)  
X.shape  
x1.size  
y_pred.shape
```

Out [32]:

(100, 1)

5. Define the classification loss function

Mean Square Error

$$L(w) = \frac{1}{n} \sum_{i=1}^n (\sigma(w^T x_i) - y_i)^2$$

Cross-Entropy

$$L(w) = \frac{1}{n} \sum_{i=1}^n (-y_i \log(\sigma(w^T x_i)) - (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

The vectorized representation is for the mean square error is as follows:

$$L(w) = \frac{1}{n} (p_w(x) - y)^T (p_w(x) - y)$$

The vectorized representation is for the cross-entropy error is as follows:

$$L(w) = \frac{1}{n} (-y^T \log(p_w(x)) - (1 - y)^T \log(1 - p_w(x)))$$

where

$$p_w(x) = \sigma(Xw) = \begin{bmatrix} \sigma(w_0 + w_1 x_{1(1)} + w_2 x_{1(2)}) \\ \sigma(w_0 + w_1 x_{2(1)} + w_2 x_{2(2)}) \\ \vdots \\ \sigma(w_0 + w_1 x_{n(1)} + w_2 x_{n(2)}) \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

You may use numpy functions `.T` and `np.log`.

In [33]:

```
def mse_loss(label, h_arr): # mean square error
    temp=h_arr-label
    L=temp*temp
    return np.mean(L)

def ce_loss(label, h_arr): # cross-entropy error
    # L=np.dot(-y.T,np.log(h_arr))-np.dot((1-y).T,np.log(1-h_arr))
    L=-label*np.log(h_arr)-(1-label)*np.log(1-h_arr)
    return np.mean(L)
```


In [43]:

```
y = data[:,2][:,None]
# temp=y-y_pred
# temp*temp
# print(np.sum(temp*temp))
print(ce_loss(y,y_pred))
sadtry = -y*np.log(y_pred)-(1-y)*np.log(1-y_pred)
sadtry
```

10.391644147841939

Out[43]:

```
array([[2.41881626e-10],
       [1.44450565e-07],
       [7.61524510e-10],
       [2.12434505e+01],
       [1.71656017e+01],
       [5.28955493e-08],
       [2.31916187e+01],
       [1.18083356e+01],
       [1.98742353e+01],
       [1.02634316e+01],
       [3.16224479e-04],
       [1.80572699e-04],
       [1.70656874e+01],
       [2.26072795e+01],
       [5.88450511e-10],
       [2.24443648e+01],
       [1.36410807e+01],
       [3.57683917e-06],
       [2.15192766e+01],
       [1.18173217e+01],
       [7.27920759e-06],
       [1.41921963e+01],
       [4.05687094e-07],
       [2.00851557e-07],
       [1.60020630e+01],
       [1.77637902e+01],
       [1.09453254e+01],
       [2.14156764e-04],
       [9.48673209e-07],
       [4.96647347e-09],
       [1.84236485e+01],
       [1.28699474e+01],
       [2.73451462e-08],
       [1.86820316e+01],
       [1.67097214e-09],
       [3.12310867e-07],
       [3.48576723e-12],
       [1.97639132e+01],
       [1.96827569e-05],
       [4.04172251e-10],
       [1.28713724e+01],
       [6.69666265e-07],
       [1.36694487e+01],
       [5.08413236e-05],
       [7.83287911e-07],
       [6.87687610e-07],
       [1.63723366e+01],
       [1.75684046e+01],
       [2.31464584e+01],
       [1.85827611e+01],
       [1.68381421e+01],
       [1.22725582e+01],
       [9.62351511e+00],
       [8.13722104e-09],
       [3.27342663e-07],
       [4.12703094e-08],
       [1.40077520e+01],
       [5.86519722e-12],
       [1.65400452e+01],
```

```
[1.85110663e+01],  
[1.96124262e+01],  
[1.27456390e-07],  
[4.89686575e-06],  
[4.51754269e-08],  
[6.68639446e-09],  
[9.51791327e-06],  
[2.54612820e+01],  
[1.91307044e-07],  
[2.03952548e+01],  
[1.55236071e+01],  
[2.07289097e-07],  
[1.92024056e+01],  
[2.20108692e+01],  
[1.85732111e+01],  
[1.92875940e+01],  
[1.44910651e+01],  
[2.29687461e+01],  
[2.01161559e+01],  
[3.89378757e-06],  
[3.29234295e-05],  
[1.50693684e+01],  
[9.65547476e+00],  
[1.65859450e+01],  
[1.61789859e+01],  
[2.01553540e+01],  
[2.02720092e+01],  
[4.32511693e-10],  
[2.05371376e+01],  
[2.14659430e+01],  
[4.48877605e-08],  
[1.60223880e+01],  
[1.84569033e+01],  
[9.48376354e-06],  
[1.95197581e+01],  
[1.00871177e+01],  
[1.13271529e+01],  
[2.31946001e+01],  
[1.38235820e+01],  
[1.74523859e+01],  
[2.04283733e+01]])
```

6. Define the gradient of the classification loss function

Given the mean square loss

$$L(w) = \frac{1}{n} \left(p_w(x) - y \right)^T \left(p_w(x) - y \right)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T \left((p_w(x) - y) \odot (p_w(x) \odot (1 - p_w(x))) \right)$$

Given the cross-entropy loss

$$L(w) = \frac{1}{n} \left(-y^T \log(p_w(x)) - (1 - y)^T \log(1 - p_w(x)) \right)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T (p_w(x) - y)$$

Implement the vectorized version of the gradient of the classification loss function

In [35]:

```
# gradient function definition(Cross Engropy)
def loss_logreg(y_pred,y):

    n = len(y)
    loss = np.dot(X.T,(y_pred-y))/n*2

    return loss

# Test loss function
y = data[:,2][:,None] # label
y_pred = f_pred(X,w) # prediction

loss = loss_logreg(y_pred,y)
loss
```

Out[35]:

```
array([[ -1.19997748],
       [-89.66073498],
       [-88.74679614]])
```

In [36]:

```
# gradient function (MeanSquareError)
def gradient_mse(y_pred,y):

    grad = np.matmul(X.T,y_pred-y)/y.size*2

    return grad
```

In [37]:

```
X.shape  
w
```

Out[37]:

```
array([[ -10.  ],  
       [  0.1 ],  
       [ -0.2 ]])
```

7. Implement the gradient descent algorithm

Vectorized implementation for the mean square loss:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T \left((p_w(x) - y) \odot (p_w(x) \odot (1 - p_w(x))) \right)$$

Vectorized implementation for the cross-entropy loss:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T (p_w(x) - y)$$

Plot the loss values $L(w^k)$ w.r.t. iteration k the number of iterations for the both loss functions.

In [52]:

```
# cross_entropy gradient descent function definition
def cross_grad_desc(X, y , w_init=np.array([0,0,0])[:,None] ,tau=1e-4, max_iter=500):

    L_iters = np.zeros([max_iter]) # record the loss values
    w_iters = np.zeros([max_iter,2]) # record the loss values
    w = w_init # initialization

    for i in range(max_iter): # loop over the iterations


        y_pred = f_pred(X,w) # linear prediction function
        #grad_f = grad_loss(y_pred,y,X) # gradient of the loss
        grad_f = loss_logreg(y_pred,y) # gradient of the loss
        w = w - tau* grad_f # update rule of gradient descent

    #     print(i)
    #     print(y_pred)
    L_iters[i] = ce_loss(y,y_pred) # save the current loss value
    #     print(L_iters[i])
    w_iters[i,:] = w[0],w[1] # save the current w value

    return w, L_iters, w_iters

# run gradient descent algorithm
start = time.time()
w_init = np.array([-10,0.1,-0.2])[:,None]
tau = 1e-4; max_iter = 1000
w, L_iters, w_iters = cross_grad_desc(X,y,w_init,tau,max_iter)

# plot
plt.figure(figsize=(14,10))
plt.plot(L_iters)
plt.xlabel('Iterations')
plt.ylabel('Loss value')
plt.show()
```



In [39]:

```
# mse gradient descent function definition
def mse_grad_desc(X, y , w_init=np.array([0,0,0])[:,None] ,tau=1e-4, max_iter=500):

    L_iters = np.zeros([max_iter]) # record the loss values
    w_iters = np.zeros([max_iter,2]) # record the loss values
    w = w_init # initialization

    for i in range(max_iter): # loop over the iterations


        y_pred = f_pred2(X,w) # linear prediction function
        #grad_f = grad_loss(y_pred,y,X) # gradient of the loss
        grad_f = gradient_mse(y_pred,y) # gradient of the loss
        w = w - tau* grad_f # update rule of gradient descent

    #     print(i)
    #     print(y_pred)
    L_iters[i] = mse_loss(y,y_pred) # save the current loss value
    #     print(L_iters[i])
    w_iters[i,:] = w[0],w[1] # save the current w value

    return w, L_iters, w_iters

# run gradient descent algorithm
start = time.time()
w_init = np.array([-10,0.1,-0.2])[:,None]
tau = 1e-4; max_iter = 1000
w, L_iters, w_iters = mse_grad_desc(X,y,w_init,tau,max_iter)

# plot
plt.figure(figsize=(14,10))
plt.plot(L_iters)
plt.xlabel('Iterations')
plt.ylabel('Loss value')
plt.show()
```



8. Plot the decision boundary

The decision boundary is defined by all points

$$x = (x_{(1)}, x_{(2)}) \quad \text{such that} \quad p_w(x) = 0.5$$

You may use numpy and matplotlib functions `np.meshgrid`, `np.linspace`, `reshape`, `contour`.

In [197]:


```

## Cross Entropy

# compute values  $p(x)$  for multiple data points  $x$ 
x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
# xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create mesh
# grid
# X2 = np.ones([np.prod(xx1.shape),3])
# X2[:,1] = xx1.reshape(-1)
# X2[:,2] = xx2.reshape(-1)
# p = f_pred(X2,w)
# p = p.reshape(50,50)
x_line=np.linspace(x1_min, x1_max)
y_line=-(w[0]+w[1]*x_line)/w[2]

# plot
plt.figure(4,figsize=(6,6))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )
plt.plot(x_line,y_line)
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Boundary", "Admitted", "Rejected"])
plt.title('Decision boundary')
plt.show()

```




In [257]:

```
## MSE

# compute values  $p(x)$  for multiple data points  $x$ 
x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
x_line=np.linspace(x1_min, x1_max)
y_line=-(w[0]+w[1]*x_line)/w[2]

# plot
plt.figure(4,figsize=(6,6))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )
plt.plot(x_line,y_line)
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Boundary", "Admitted", "Rejected"])
plt.title('Decision boundary')
plt.show()
```



In [192]:

```
x_line=np.linspace(x1_min, x1_max)
y_line=-(w[0]+w[1]*x_line)/w[2]
```

9. Comparison with Scikit-learn logistic regression algorithm with the gradient descent with the cross-entropy loss

You may use scikit-learn function `LogisticRegression(C=1e6)` .

In [58]:

```

# run logistic regression with scikit-learn
start = time.time()
logreg_sklearn = LogisticRegression(C=1e6) # scikit-learn logistic regression
logreg_sklearn.fit(data[:,[0,1]],data[:,2]) # learn the model parameters

# compute loss value
b=logreg_sklearn.intercept_
w_sklearn = logreg_sklearn.coef_

# plot
plt.figure(4,figsize=(6,6))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )


plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')

x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
x_line=np.linspace(x1_min, x1_max)
y_line=-(b[0]+w_sklearn[0][0]*x_line)/w_sklearn[0][1]

plt.plot(x_line,y_line)
y_line2=-(w[0]+w[1]*x_line)/w[2]
plt.plot(x_line,y_line2)

plt.title('Decision boundary (black with gradient descent and magenta with scikit-learn)')
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Boundary", "Admitted", "Rejected"])
plt.title('Decision boundary')
plt.show()

```



In [48]:

```
start = time.time()
logreg_sklearn = LogisticRegression(C=1e6) # scikit-learn logistic regression
logreg_sklearn.fit(data[:,[0,1]],data[:,2]) # learn the model parameters
# logreg_sklearn.fit(X,data[:,2])
logreg_sklearn.coef_
logreg_sklearn.predict(data[:,[0,1]])
w_sklearn = logreg_sklearn.coef_
logreg_sklearn.intercept_
```

Out [48]:

```
array([-25.16138556])
```

In [53]:

```
w
```

Out [53]:

```
array([[ -10.00125246],
       [  0.08547625],
       [  0.07922129]])
```

10. Plot the probability map

In [248]:

```

#cross entropy
num_a = 110
grid_x1 = np.linspace(20,110,num_a)
grid_x2 = np.linspace(20,110,num_a)

score_x1, score_x2 = np.meshgrid(grid_x1,grid_x2)

X_temp = np.ones([np.prod(score_x1.shape),3])
X_temp[:,1]=score_x1.reshape(-1)
X_temp[:,2]=score_x2.reshape(-1)
Z = f_pred(X_temp,w).reshape(110,110)
# xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create mesh
grid
# X2 = np.ones([np.prod(xx1.shape),3])
# X2[:,1] = xx1.reshape(-1)
# X2[:,2] = xx2.reshape(-1)
# p = f_pred(X2,w)
# p = p.reshape(50,50)
# for i in range(len(score_x1)):
#     for j in range(len(score_x2)):

#         predict_prob = sigmoid( )


#         Z[j, i] = predict_prob

# actual plotting example
fig = plt.figure(figsize=(10,10))
ax = fig.add_subplot(111)
ax.tick_params(length=6, width=2, grid_alpha=0.5)
ax.set_xlabel('Exam grade 1')
ax.set_ylabel('Exam grade 2')

ax.set_xlim(20, 110)
ax.set_ylim(20, 110)

cf = ax.contourf(score_x1,score_x2,Z,cmap="bwr",alpha=0.5)
ax.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+",label="Admitted")
ax.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o",label="Rejected")
cbar = fig.colorbar(cf)
cbar.update_ticks()
plt.legend()
plt.show()

```



In [40]:

```
##MSE
num_a = 110
grid_x1 = np.linspace(20,110,num_a)
grid_x2 = np.linspace(20,110,num_a)


score_x1, score_x2 = np.meshgrid(grid_x1,grid_x2)

X_temp = np.ones([np.prod(score_x1.shape),3])
X_temp[:,1]=score_x1.reshape(-1)
X_temp[:,2]=score_x2.reshape(-1)
Z = f_pred2(X_temp,w).reshape(110,110)

fig = plt.figure(figsize=(10,10))
ax = fig.add_subplot(111)
ax.tick_params(length=6, width=2, grid_alpha=0.5)
ax.set_xlabel('Exam grade 1')
ax.set_ylabel('Exam grade 2')

ax.set_xlim(20, 110)
ax.set_ylim(20, 110)

cf = ax.contourf(score_x1,score_x2,Z,cmap="bwr",alpha=0.5)
ax.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+",label="Admitted")
ax.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o",label="Rejected" )
cbar = fig.colorbar(cf)
cbar.update_ticks()
plt.legend()
plt.show()
```



In [210]:

```

num_a = 110
grid_x1 = np.linspace(20,110,num_a)
grid_x2 = np.linspace(20,110,num_a)

score_x1, score_x2 = np.meshgrid(grid_x1,grid_x2)

Z = np.zeros(score_x1.shape)
Z.shape
X_temp = np.ones([np.prod(score_x1.shape),3])
X_temp[:,1]=score_x1.reshape(-1)
X_temp[:,2]=score_x2.reshape(-1)
Z = f_pred(X_temp,w).reshape(110,110)
Z.shape
Z

```

Out [210]:

```

array([[0.00122047, 0.0013096 , 0.00140523, ..., 0.69939273, 0.71401906,
        0.72821029],
       [0.00130286, 0.001398 , 0.00150008, ..., 0.7129633 , 0.7271869 ,
        0.74096176],
       [0.00139081, 0.00149236, 0.00160132, ..., 0.7261611 , 0.73996924,
        0.75331763],
       ...,
       [0.57243439, 0.58961332, 0.60657641, ..., 0.99960786, 0.99963457,
        0.99965947],
       [0.58836306, 0.60534323, 0.62207289, ..., 0.99963268, 0.9996577 ,
        0.99968102],
       [0.60410872, 0.62085792, 0.63732393, ..., 0.99965593, 0.99967937,
        0.99970121]])


```

Output results

1. Plot the dataset in 2D cartesian coordinate system (1pt)

In [18]:


```
plt.figure(figsize=(8,8))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )
plt.title('Training data')
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Admitted", "Rejected"])
plt.show()
```



2. Plot the sigmoid function (1pt)

In [20]:


```
plt.figure(2)
plt.plot(x_values, sigmoid(x_values))
plt.title("Sigmoid function")
plt.grid(True)
```



3. Plot the loss curve in the course of gradient descent using the mean square error (2pt)

In [256]:


```
plt.figure(figsize=(8,6))
plt.plot(L_iters)
plt.xlabel('Iterations')
plt.ylabel('Loss value')
plt.show()
```



4. Plot the loss curve in the course of gradient descent using the cross-entropy error (2pt)

In [174]:


```
# plot
plt.figure(figsize=(8,6))
plt.plot(L_iters)
plt.xlabel('Iterations')
plt.ylabel('Loss value')
plt.show()
```



5. Plot the decision boundary using the mean square error (2pt)

In [258]:


```
plt.figure(4,figsize=(6,6))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )
plt.plot(x_line,y_line)
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Boundary", "Admitted", "Rejected"])
plt.title('Decision boundary')
plt.show()
```



6. Plot the decision boundary using the cross-entropy error (2pt)

In [198]:

```
plt.figure(4,figsize=(6,6))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )
plt.plot(x_line,y_line)
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Boundary", "Admitted", "Rejected"])
plt.title('Decision boundary')
plt.show()
```



7. Plot the decision boundary using the Scikit-learn logistic regression algorithm (2pt)

In [57]:


```
plt.figure(4,figsize=(6,6))
plt.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+")
plt.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o" )

plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')

x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
x_line=np.linspace(x1_min, x1_max)
y_line=-(b[0]+w_sklearn[0][0]*x_line)/w_sklearn[0][1]

plt.plot(x_line,y_line)

plt.title('Decision boundary (black with gradient descent and magenta with scikit-learn)')
plt.xlabel('Exam grade 1')
plt.ylabel('Exam grade 2')
plt.legend(["Boundary", "Admitted", "Rejected"])
plt.title('Decision boundary')
plt.show()
```



8. Plot the probability map using the mean square error (2pt)

In [42]:


```

fig = plt.figure(figsize=(10,10))
ax = fig.add_subplot(111)
ax.tick_params(length=6, width=2, grid_alpha=0.5)
ax.set_xlabel('Exam grade 1')
ax.set_ylabel('Exam grade 2')

ax.set_xlim(20, 110)
ax.set_ylim(20, 110)

cf = ax.contourf(score_x1,score_x2,Z,cmap="bwr",alpha=0.5)
ax.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+",label="Admitted")
ax.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o",label="Rejected" )
cbar = fig.colorbar(cf)
cbar.update_ticks()
plt.legend()
plt.show()

```




9. Plot the probability map using the cross-entropy error (2pt)

In [249]:

```
fig = plt.figure(figsize=(10,10))
ax = fig.add_subplot(111)
ax.tick_params(length=6, width=2, grid_alpha=0.5)
ax.set_xlabel('Exam grade 1')
ax.set_ylabel('Exam grade 2')

ax.set_xlim(20, 110)
ax.set_ylim(20, 110)

cf = ax.contourf(score_x1,score_x2,Z,cmap="bwr",alpha=0.5)
ax.scatter(x=x1[idx_admit],y=x2[idx_admit],c="red",marker="+",label="Admitted")
ax.scatter(x1[idx_rejec],x2[idx_rejec],c="blue",marker="o",label="Rejected" )
cbar = fig.colorbar(cf)
cbar.update_ticks()
plt.legend()
plt.show()
```



In []: